


How Gemini is wired

The frontend never calls Gemini. Slack never calls Gemini. Firestore never calls Gemini. One backend process — FastAPI on Cloud Run — is the only client of `google.genai`. Sage arrows are the only Gemini path. The dashed line is the missing Gemini→Firestore session.



Who talks to whom

Gemini is a backend RPC. It has no session to the Replay UI, Slack, or Firestore.

FROM	TO	HOW
Replay UI	FastAPI	fetch POST /api/run — no Gemini key in the browser
Slack	FastAPI	/slack/events · /slack/commands · Bolt ACK then pipeline
MCP	FastAPI	POST /mcp/query_decisions
FastAPI	Gemini 3.5	gemini_client.py · generate_content + embed_content
FastAPI	Firestore	graph().upsert · labels, status, permalinks, edges
FastAPI	Slack search	slack_store · live snippets, discarded after the request
Gemini	Replay UI	Never — the browser only reads the JSON card from FastAPI
Gemini	Firestore	Never — the pipeline writes derived metadata after adjudication

Request path

- 1 Replay posts POST /api/run. Slack posts /slack/commands or /slack/events. MCP posts /mcp/query_decisions.
- 2 Prefilter drops chatter with no model. Gate uses Flash-Lite. About 95% of watcher traffic exits here.
- 3 Flash-Lite writes cross-vocabulary probes. Retrieve fans out to live Slack search and the Firestore label graph.
- 4 embed() ranks candidates in memory, then drops the vectors. Flash adjudicates same / current / reversed / concurrent.
- 5 FastAPI writes derived labels to Firestore and returns a card to Replay or Block Kit to Slack. Gemini never sees the database.

What crosses the database

STORED	NEVER STORED
+ derived labels	× Slack message text
+ status	× Slack user IDs
+ confidence	× content embeddings
+ permalinks	× a persistent vector index
+ edges	

Models

Gate / probes	gemini-3.5-flash-lite
Adjudicate	gemini-3.5-flash
Transient rank	gemini-embedding-001
Framework	Google ADK sequential pipeline